

Assignment 7: Time Series Analysis

Abhay V Rao

OVERVIEW

This exercise accompanies the lessons in Environmental Data Analytics on time series analysis.

Directions

1. Change “Student Name” on line 3 (above) with your name.
2. Work through the steps, **creating code and output** that fulfill each instruction.
3. Be sure to **answer the questions** in this assignment document.
4. When you have completed the assignment, **Knit** the text and code into a single PDF file.
5. After Knitting, submit the completed exercise (PDF file) to the dropbox in Sakai. Add your last name into the file name (e.g., “Fay_A07_TimeSeries.Rmd”) prior to submission.

The completed exercise is due on Monday, March 14 at 7:00 pm.

Set up

1. Set up your session:
 - Check your working directory
 - Load the tidyverse, lubridate, zoo, and trend packages
 - Set your ggplot theme

```
#1
getwd()

## [1] "C:/Users/av241/Documents/Environmental_Data_Analytics_2022/Assignments"

library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --
## v ggplot2 3.3.5      v purrr  0.3.4
## v tibble  3.1.6      v dplyr  1.0.7
## v tidyr   1.1.4      v stringr 1.4.0
## v readr   2.1.1      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

library(lubridate)

##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
##
##     date, intersect, setdiff, union
```

```
library(zoo)

##
## Attaching package: 'zoo'
## The following objects are masked from 'package:base':
##
##      as.Date, as.Date.numeric

library(trend)

mytheme <- theme_gray(base_size = 10) +
  theme(axis.text = element_text(color = "darkgrey"),
        legend.position = "bottom")
theme_set(mytheme)
```

2. Import the ten datasets from the Ozone_TimeSeries folder in the Raw data folder. These contain ozone concentrations at Garinger High School in North Carolina from 2010-2019 (the EPA air database only allows downloads for one year at a time). Import these either individually or in bulk and then combine them into a single dataframe named **GaringerOzone** of 3589 observation and 20 variables.

```
#2
Ozone1 <- read.csv("../Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2010_raw.csv",
                  stringsAsFactors = TRUE)
Ozone2 <- read.csv("../Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2011_raw.csv",
                  stringsAsFactors = TRUE)
Ozone3 <- read.csv("../Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2012_raw.csv",
                  stringsAsFactors = TRUE)
Ozone4 <- read.csv("../Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2013_raw.csv",
                  stringsAsFactors = TRUE)
Ozone5 <- read.csv("../Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2014_raw.csv",
                  stringsAsFactors = TRUE)
Ozone6 <- read.csv("../Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2015_raw.csv",
                  stringsAsFactors = TRUE)
Ozone7 <- read.csv("../Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2016_raw.csv",
                  stringsAsFactors = TRUE)
Ozone8 <- read.csv("../Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2017_raw.csv",
                  stringsAsFactors = TRUE)
Ozone9 <- read.csv("../Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2018_raw.csv",
                  stringsAsFactors = TRUE)
Ozone10 <- read.csv("../Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2019_raw.csv",
                   stringsAsFactors = TRUE)

GaringerOzone <- rbind(Ozone1,Ozone2,Ozone3,Ozone4,Ozone5,Ozone6,Ozone7,
                      Ozone8,Ozone9,Ozone10)

dim(GaringerOzone) #Dimensions are 3589*20

## [1] 3589    20
```

Wrangle

3. Set your date column as a date class.
4. Wrangle your dataset so that it only contains the columns Date, Daily.Max.8.hour.Ozone.Concentration, and DAILY_AQI_VALUE.

5. Notice there are a few days in each year that are missing ozone concentrations. We want to generate a daily dataset, so we will need to fill in any missing days with NA. Create a new data frame that contains a sequence of dates from 2010-01-01 to 2019-12-31 (hint: `as.data.frame(seq())`). Call this new data frame `Days`. Rename the column name in `Days` to “Date”.
6. Use a `left_join` to combine the data frames. Specify the correct order of data frames within this function so that the final dimensions are 3652 rows and 3 columns. Call your combined data frame `GaringerOzone`.

```
# 3
GaringerOzone$Date <- mdy(GaringerOzone$Date)
class(GaringerOzone$Date)

## [1] "Date"

# 4
GaringerOzoneAnalysis <- GaringerOzone %>%
  select(Date, Daily.Max.8.hour.Ozone.Concentration, DAILY_AQI_VALUE)

# 5
Days <- as.data.frame(seq(as.Date('2010/01/01'), as.Date('2019/12/31'), by="1 days"))
colnames(Days)[1] <- "Date"

# 6
GaringerOzone <- left_join(Days,GaringerOzoneAnalysis)

## Joining, by = "Date"
dim(GaringerOzone)


## [1] 3652    3
#'Days' column replaced, dimensions are now 3652*3
```

Visualize

7. Create a line plot depicting ozone concentrations over time. In this case, we will plot actual concentrations in ppm, not AQI values. Format your axes accordingly. Add a smoothed line showing any linear trend of your data. Does your plot suggest a trend in ozone concentration over time?

```
#7
ggplot(GaringerOzone, aes(x = Date, y = Daily.Max.8.hour.Ozone.Concentration)) +
  geom_line() +
  geom_smooth(method= "lm") +
  labs(x = "Date", y = "PPM")

## `geom_smooth()` using formula 'y ~ x'
## Warning: Removed 63 rows containing non-finite values (stat_smooth).
```



There is a slight downward trend in ozone concentrations over time.

Answer:

Time Series Analysis

Study question: Have ozone concentrations changed over the 2010s at this station?

8. Use a linear interpolation to fill in missing daily data for ozone concentration. Why didn't we use a piecewise constant or spline interpolation?

```
#8
GaringerOzoneClean <-
  GaringerOzone %>%
  mutate( Ozone.Concentration = zoo::na.approx(Daily.Max.8.hour.Ozone.Concentration) )
```

#Interpolating to cover NA gaps

Answer:

We use linear interpolation because the gaps in data are small. A spline/ quadratic interpolation is not relevant for this data based on the linear trendline in qn 7.

9. Create a new data frame called `GaringerOzone.monthly` that contains aggregated data: mean ozone concentrations for each month. In your pipe, you will need to first add columns for year and month to form the groupings. In a separate line of code, create a new Date column with each month-year combination being set as the first day of the month (this is for graphing purposes only)

#9

```

GaringerOzone.monthly <- GaringerOzoneClean %>%
  mutate(month=month(GaringerOzoneClean$Date)) %>%
  mutate(year=year(GaringerOzoneClean$Date)) %>%
  mutate(month.year = my(paste0(month,"-",year))) %>%
  group_by(month.year) %>%
  summarize(MeanOzone = mean(Ozone.Concentration))

```

*# Create month, year and concatenate; then group by concatenated field. PPMs are
then summarized by mean.*

10. Generate two time series objects. Name the first `GaringerOzone.daily.ts` and base it on the dataframe of daily observations. Name the second `GaringerOzone.monthly.ts` and base it on the monthly average ozone values. Be sure that each specifies the correct start and end dates and the frequency of the time series.

#10

```

GaringerOzone.daily.ts <- ts(GaringerOzoneClean$Ozone.Concentration, start(2010,1),frequency= 365)
GaringerOzone.monthly.ts <- ts(GaringerOzone.monthly$MeanOzone, start(2010,1),frequency=12)

```


11. Decompose the daily and the monthly time series objects and plot the components using the `plot()` function.

#11


```

GaringerDailyDecomp <- stl(GaringerOzone.daily.ts, s.window= "periodic")
plot(GaringerDailyDecomp)

```



```
GaringerMonthlyDecomp <- stl(GaringerOzone.monthly.ts,s.window = "periodic")
plot(GaringerMonthlyDecomp)
```



12. Run a monotonic trend analysis for the monthly Ozone series. In this case the seasonal Mann-Kendall is most appropriate; why is this?

```
#12
Monthly.Ozone.Garinger.Trend1 <- Kendall::SeasonalMannKendall(GaringerOzone.monthly.ts)
Monthly.Ozone.Garinger.Trend1
```

```
## tau = -0.143, 2-sided pvalue =0.046724
```

```
summary(Monthly.Ozone.Garinger.Trend1)
```

```
## Score = -77 , Var(Score) = 1499
```

```
## denominator = 539.4972
```

```
## tau = -0.143, 2-sided pvalue =0.046724
```

```
Monthly.Ozone.Garinger.Trend2 <- trend::smk.test(GaringerOzone.monthly.ts)
```

```
Monthly.Ozone.Garinger.Trend2
```

```
##
```

```
## Seasonal Mann-Kendall trend test (Hirsch-Slack test)
```

```
##
```

```
## data: GaringerOzone.monthly.ts
```

```
## z = -1.963, p-value = 0.04965
```

```
## alternative hypothesis: true S is not equal to 0
```

```
## sample estimates:
##      S varS
##    -77 1499
```

```
summary(Monthly.Ozone.Garinger.Trend2)
```

```
##
## Seasonal Mann-Kendall trend test (Hirsch-Slack test)
##
## data: GaringerOzone.monthly.ts
## alternative hypothesis: two.sided
##
## Statistics for individual seasons
##
## H0
##
##      S varS      tau      z Pr(>|z|)
## Season 1:  S = 0   -1  125 -0.022  0.000  1.00000
## Season 2:  S = 0   -4  124 -0.090 -0.269  0.78762
## Season 3:  S = 0  -17  125 -0.378 -1.431  0.15241
## Season 4:  S = 0 -15  125 -0.333 -1.252  0.21050
## Season 5:  S = 0 -17  125 -0.378 -1.431  0.15241
## Season 6:  S = 0 -11  125 -0.244 -0.894  0.37109
## Season 7:  S = 0   -7  125 -0.156 -0.537  0.59151
## Season 8:  S = 0   -5  125 -0.111 -0.358  0.72051
## Season 9:  S = 0 -13  125 -0.289 -1.073  0.28313
## Season 10: S = 0 -13  125 -0.289 -1.073  0.28313
## Season 11: S = 0  11  125  0.244  0.894  0.37109
## Season 12: S = 0  15  125  0.333  1.252  0.21050
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


Answer: There is visible seasonality to the data upon decomposition so the seasonal Mann Kendall is the most appropriate. The trend analysis indicates, based on the p-value of 0.04965, that we can reject the null hypothesis that the data is stationary.

13. Create a plot depicting mean monthly ozone concentrations over time, with both a `geom_point` and a `geom_line` layer. Edit your axis labels accordingly.

```
# 13
```

```
Plot.Monthly.Ozone.Garinger <- ggplot(GaringerOzone.monthly, aes(x=month.year, y=MeanOzone))+
  geom_point()+
  geom_line()+
  xlab("Date")+
  ylab("Mean Ozone Concentration")+
  geom_smooth(method=lm, se=FALSE)
print(Plot.Monthly.Ozone.Garinger)
```

```
## `geom_smooth()` using formula 'y ~ x'
```



14. To accompany your graph, summarize your results in context of the research question. Include output from the statistical test in parentheses at the end of your sentence. Feel free to use multiple sentences in your interpretation.


Answer: There is a negative slope over the ten year timeframe, indicating reducing ozone concentrations. The overall seasonal Mann Kendall p-value of < 0.05 confirms a monotonic trend.

15. Subtract the seasonal component from the `GaringerOzone.monthly.ts`. Hint: Look at how we extracted the series components for the `EnoDischarge` on the lesson Rmd file.
16. Run the Mann Kendall test on the non-seasonal Ozone monthly series. Compare the results with the ones obtained with the Seasonal Mann Kendall on the complete series.

```
#15
# Creating data frame to see how the trend maps onto the original data
Monthly.Ozone.Components <- as.data.frame(GaringerMonthlyDecomp$time.series[,2:3])
Monthly.Ozone.Components.2 <- mutate(Monthly.Ozone.Components, PPM = GaringerOzone.monthly$MeanOzone,
                                     Date = GaringerOzone.monthly$month.year)%>%
  select(trend, remainder, PPM, Date)
# adding on the mean PPM and date columns to the decomposed TS

#Creating new TS
TS.Ozone.2 <- ts(Monthly.Ozone.Components.2$PPM, start=c(2010,1), frequency= 12)

#decomposing
GaringerMonthlyDecomp.2 <- stl(TS.Ozone.2,s.window = "periodic")
plot(GaringerMonthlyDecomp.2)
```

#16

```
Nonseasonal.Ozone.Garinger.Trend1 <- Kendall::MannKendall(TS.Ozone.2)
Nonseasonal.Ozone.Garinger.Trend1
```

```
## tau = -0.0594, 2-sided pvalue =0.33732
```

```
summary(Nonseasonal.Ozone.Garinger.Trend1)
```

```
## Score = -424 , Var(Score) = 194364.7
```

```
## denominator = 7139
```

```
## tau = -0.0594, 2-sided pvalue =0.33732
```

Answer: The Mann Kendall Test has a p-value of 0.33732, which accepts the null hypothesis, unlike the Seasonal Test.